

CS6005 Deep Learning Techniques
Natural Language Processing Project
Spam V/S Ham SMS Classification



Name	Gokul. S
Roll Number	2018103026
Batch	P
Semester	6 (RUSA)
Date of Submission	21. 05. 2021
Github Profile	www.github.com/gokul-sunilkumar

Problem Statement: To classify an SMS message as either a spam/ham using text analytics, natural language processing and machine learning

Dataset: SMS Spam Collection Data Set

Description: The dataset consists of more than 5000 SMS phone messages belonging to two classes namely spam, ham.

Machine Learning Repository

Center for Machine Learning and Intelligent Systems

SMS Spam Collection Data Set

Download: [Data Folder](#), [Data Set Description](#)

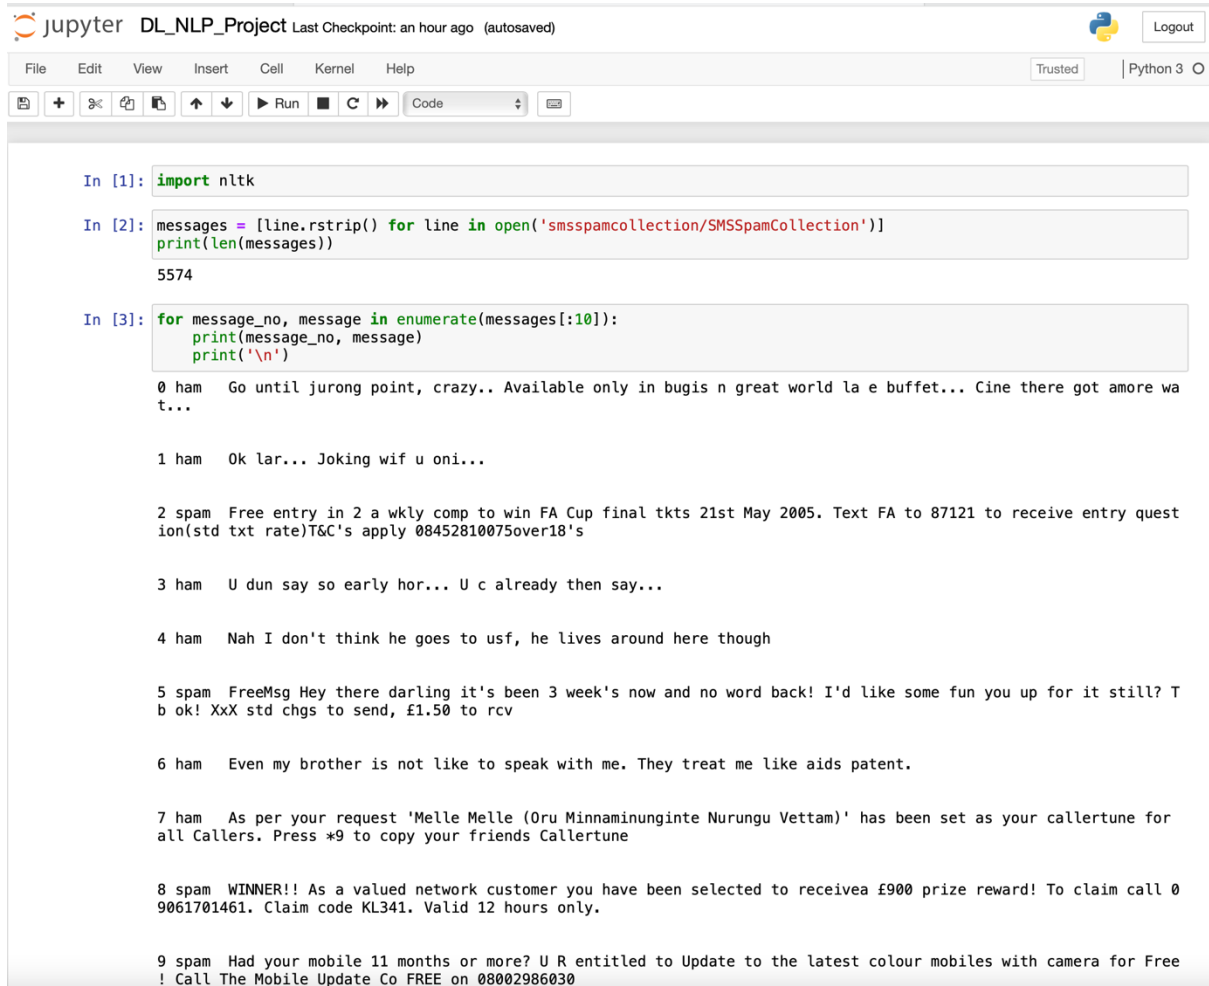
Abstract: The SMS Spam Collection is a public set of SMS labeled messages that have been collected for mobile phone spam research.

Data Set Characteristics:	Multivariate, Text, Domain-Theory	Number of Instances:	5574	Area:	Computer
Attribute Characteristics:	Real	Number of Attributes:	N/A	Date Donated	2012-06-22
Associated Tasks:	Classification, Clustering	Missing Values?	N/A	Number of Web Hits:	358386

URL: (<https://archive.ics.uci.edu/ml/datasets/sms+spam+collection>) → UCI Machine Learning Repository

Code/ Execution Snapshots:

Importing packages and checking the structure, shape and contents of the dataset.



The image shows a Jupyter Notebook interface with the title 'DL_NLP_Project' and a status bar indicating 'Last Checkpoint: an hour ago (autosaved)'. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Help) and a toolbar with icons for file operations, running, and code execution. The notebook contains three code cells:

```
In [1]: import nltk
```

```
In [2]: messages = [line.rstrip() for line in open('smsspamcollection/SMSSpamCollection')]
print(len(messages))
5574
```

```
In [3]: for message_no, message in enumerate(messages[:10]):
        print(message_no, message)
        print('\n')
```

The output of the third cell shows the first 10 messages from the dataset, including ham and spam messages:

```
0 ham  Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wa
t...

1 ham  Ok lar... Joking wif u oni...

2 spam Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry quest
ion(std txt rate)T&C's apply 08452810075over18's

3 ham  U dun say so early hor... U c already then say...

4 ham  Nah I don't think he goes to usf, he lives around here though

5 spam FreeMsg Hey there darling it's been 3 week's now and no word back! I'd like some fun you up for it still? T
b ok! XxX std chgs to send, £1.50 to rcv


6 ham  Even my brother is not like to speak with me. They treat me like aids patent.


7 ham  As per your request 'Melle Melle (Oru Minnaminunginte Nurungu Vettam)' has been set as your callertune for
all Callers. Press *9 to copy your friends Callertune

8 spam WINNER!! As a valued network customer you have been selected to receivea £900 prize reward! To claim call 0
9061701461. Claim code KL341. Valid 12 hours only.

9 spam Had your mobile 11 months or more? U R entitled to Update to the latest colour mobiles with camera for Free
! Call The Mobile Update Co FREE on 08002986030
```

Read the csv data file and create a Pandas Dataframe. Check how balanced the dataset is.

Jupyter DL_NLP_Project Last Checkpoint: an hour ago (autosaved)  Logout

File Edit View Insert Cell Kernel Help Trusted Python 3 

In [4]: `import pandas as pd`

In [5]: `messages = pd.read_csv('smsspamcollection/SMSSpamCollection', sep='\t',
names=["label", "message"])
messages.head()`

Out[5]:

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

In [6]: `messages.describe()`

Out[6]:

	label	message
count	5572	5572
unique	2	5169
top	ham	Sorry, I'll call later
freq	4825	30

In [7]: `messages.groupby('label').describe()`

Out[7]:

	count	unique	top	freq
label				
ham	4825	4516	Sorry, I'll call later	30
spam	747	653	Please call our customer service representativ...	4

Create an additional meaningful column called length and perform Exploratory Data Analysis (EDA)

Jupyter DL_NLP_Project Last Checkpoint: an hour ago (autosaved) Python 3 Logout

File Edit View Insert Cell Kernel Help Trusted

In [8]: `messages['length'] = messages['message'].apply(len)`
`messages.head()`

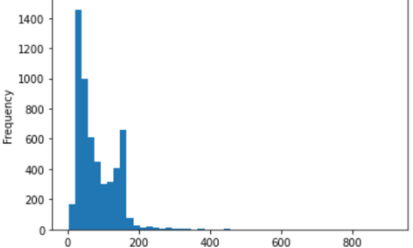
Out[8]:

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

In [9]: `import matplotlib.pyplot as plt`
`import seaborn as sns`
`%matplotlib inline`

In [10]: `messages['length'].plot(bins=50, kind='hist')`

Out[10]: `<AxesSubplot:ylabel='Frequency'>`



In [11]: `messages.length.describe()`

Out[11]:

count	5572.000000
mean	80.489950
std	59.942907
min	2.000000
25%	36.000000
50%	62.000000
75%	122.000000
max	910.000000
Name: length, dtype: float64	

Remove punctuations and stopwords from the input messages

Jupyter DL_NLP_Project Last Checkpoint: an hour ago (autosaved) Python 3 Logout

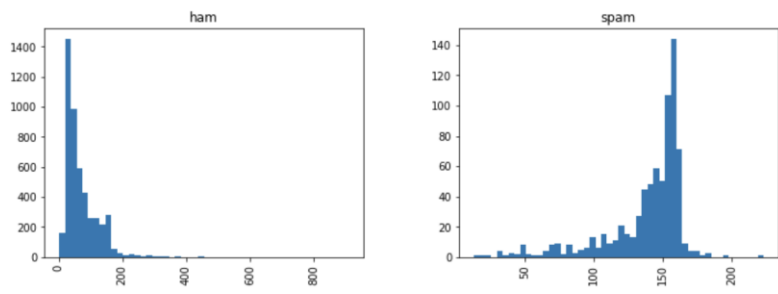
File Edit View Insert Cell Kernel Help Trusted Python 3

In [12]: `messages[messages['length'] == 910]['message'].iloc[0]`

Out[12]: "For me the love should start with attraction.i should feel that I need her every time around me.she should be the first thing which comes in my thoughts.I would start the day and end it with her.she should be there every time I dream.love will be then when my every breath has her name.my life should happen around her.my life will be named t o her.I would cry for her.will give all my happiness and take all her sorrows.I will be ready to fight with anyone for her.I will be in love when I will be doing the craziest things for her.love will be when I don't have to prove anyone that my girl is the most beautiful lady on the whole planet.I will always be singing praises for her.love will be when I start up making chicken curry and end up making sambar.life will be the most beautiful then.will g et every morning and thank god for the day because she is with me.I would like to say a lot..will tell later.."

In [13]: `messages.hist(column='length', by='label', bins=50, figsize=(12,4))`

Out[13]: `array([<AxesSubplot:title={'center':'ham'}>, <AxesSubplot:title={'center':'spam'}>], dtype=object)`



In [14]: `import string`
`mess = 'Sample message! Notice: it has punctuation.'`
`# Check characters to see if they are in punctuation`
`nopunc = [char for char in mess if char not in string.punctuation]`
`# Join the characters again to form the string.`
`nopunc = ''.join(nopunc)`

In [15]: `from nltk.corpus import stopwords`
`stopwords.words('english')[0:10] # Show some stop words`

Out[15]: `['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're"]`

In [16]: `nopunc.split()`

Out[16]: `['Sample', 'message', 'Notice', 'it', 'has', 'punctuation']`

Pre-process the dataframe using the above strategies and update it in place.

Jupyter DL_NLP_Project Last Checkpoint: an hour ago (autosaved) Python 3 Logout

File Edit View Insert Cell Kernel Help

In [17]: `# Now just remove any stopwords
clean_mess = [word for word in nopunc.split() if word.lower() not in stopwords.words('english')]`

In [18]: `clean_mess`

Out[18]: `['Sample', 'message', 'Notice', 'punctuation']`

In [19]: `def text_process(mess):
 nopunc = [char for char in mess if char not in string.punctuation]
 nopunc = ''.join(nopunc)
 return [word for word in nopunc.split() if word.lower() not in stopwords.words('english')]`

In [20]: `messages.head()`

Out[20]:

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

In [21]: `# Check to make sure its working
messages['message'].head(5).apply(text_process)`

Out[21]:


```
0    [Go, jurong, point, crazy, Available, bugis, n...  
1    [Ok, lar, Joking, wif, u, oni]  
2    [Free, entry, 2, wkly, comp, win, FA, Cup, fin...  
3    [U, dun, say, early, hor, U, c, already, say]  
4    [Nah, dont, think, goes, usf, lives, around, t...  
Name: message, dtype: object
```

In [22]: `# Show original dataframe
messages.head()`

Out[22]:

	label	message	length
0	ham	Go until jurong point, crazy.. Available only ...	111
1	ham	Ok lar... Joking wif u oni...	29
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	155
3	ham	U dun say so early hor... U c already then say...	49
4	ham	Nah I don't think he goes to usf, he lives aro...	61

Convert cleaned messages to their bag of words form using Sklearn's CountVectorizer

```
Jupyter DL_NLP_Project Last Checkpoint: an hour ago (autosaved)  Logout

File Edit View Insert Cell Kernel Help Trusted Python 3

In [23]: from sklearn.feature_extraction.text import CountVectorizer

In [24]: # Might take awhile...
bow_transformer = CountVectorizer(analyzer=text_process).fit(messages['message'])

# Print total number of vocab words
print(len(bow_transformer.vocabulary_))

11425

In [25]: message4 = messages['message'][3]
print(message4)

U dun say so early hor... U c already then say...

In [26]: bow4 = bow_transformer.transform([message4])
print(bow4)
print(bow4.shape)

(0, 4068) 2
(0, 4629) 1
(0, 5261) 1
(0, 6204) 1
(0, 6222) 1
(0, 7186) 1
(0, 9554) 2
(1, 11425)

In [27]: print(bow_transformer.get_feature_names()[4068])
print(bow_transformer.get_feature_names()[9554])

U
say

In [28]: messages_bow = bow_transformer.transform(messages['message'])


In [29]: print('Shape of Sparse Matrix: ', messages_bow.shape)
print('Amount of Non-Zero occurrences: ', messages_bow.nnz)

Shape of Sparse Matrix: (5572, 11425)
Amount of Non-Zero occurrences: 50548

In [30]: sparsity = (100.0 * messages_bow.nnz / (messages_bow.shape[0] * messages_bow.shape[1]))
print('sparsity: {}'.format(round(sparsity)))

sparsity: 0
```


Use Sklearn's learning_curve function to plot the accuracy curve for different train-test split sizes (considering 50:50, 60:40, 70:30, 80:20) with a cross validation K-Fold value of 5 and the average of the five subordinate accuracies is the overall accuracy for that train size.

```
Jupyter DL_NLP_Project Last Checkpoint: 19 hours ago (unsaved changes)  Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

In [35]: #Finding the appropriate train-test split and preventing overfitting(Using K-Fold Cross Validation on accuracy)

In [45]: from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import learning_curve

train_sizes = [1500, 2250, 2750, 3300, 3850, 4350, 4450]
t_sizes, train_scores, validation_scores = learning_curve(
    estimator=Pipeline([
        ('bow', CountVectorizer(analyzer=text_process)), # strings to token integer counts
        ('tfidf', TfidfTransformer()), # integer counts to weighted TF-IDF scores
        ('classifier', MultinomialNB()), # train on TF-IDF vectors w/ Naive Bayes classifier
    ])
    , X=messages["message"], y=messages["label"], train_sizes=train_sizes, cv=5, scoring="accuracy")

In [47]: train_scores
Out[47]: array([[0.964, 0.96866667, 0.96466667, 0.96466667, 0.96466667],
 [0.96977778, 0.97022222, 0.97422222, 0.97333333, 0.97333333],
 [0.97054545, 0.972, 0.976, 0.976, 0.976],
 [0.97484848, 0.97454545, 0.97787879, 0.97484848, 0.97484848],
 [0.97636364, 0.97636364, 0.97818182, 0.9774026, 0.97506494],
 [0.97908046, 0.97770115, 0.97908046, 0.97862069, 0.97747126],
 [0.97932584, 0.97820225, 0.97955056, 0.97865169, 0.97752809]])

In [48]: validation_scores
Out[48]: array([[0.93273543, 0.93273543, 0.93536804, 0.92908438, 0.93716338],
 [0.9470852, 0.94170404, 0.94344704, 0.93985637, 0.94883303],
 [0.95426009, 0.95067265, 0.9524237, 0.94614004, 0.95421903],
 [0.9632287, 0.95515695, 0.95780969, 0.95062837, 0.95601436],
 [0.96502242, 0.95695067, 0.95870736, 0.95332136, 0.95960503],
 [0.96591928, 0.95695067, 0.96050269, 0.95332136, 0.96229803],
 [0.96591928, 0.95695067, 0.95960503, 0.95332136, 0.96319569]])

In [49]: train_scores_mean = train_scores.mean(axis = 1)
validation_scores_mean = validation_scores.mean(axis = 1)
print('Mean training scores\n\n', pd.Series(train_scores_mean, index = train_sizes))
print('\n', '-' * 20) # separator
print('\nMean validation scores\n\n', pd.Series(validation_scores_mean, index = train_sizes))

Mean training scores
1500 0.965333
2250 0.972178
2750 0.974109
3300 0.975394
3850 0.976675
4350 0.978391
4450 0.978652
dtype: float64

-----

Mean validation scores
1500 0.933417
2250 0.944185
2750 0.951543
3300 0.956568
3850 0.958721
4350 0.959798
4450 0.959798
dtype: float64
```

Use Sklearn's learning_curve function to plot the log_loss_curve for different train-test split sizes (considering 50:50, 60:40, 70:30, 80:20) with a cross validation K-Fold value of 5 and the average of the five subordinate accuracies is the overall log loss for that train size.

Scoring parameter here is given as "neg_log_loss" as learning_curve function tries to maximise the scoring parameter.

```
jupyter DL_NLP_Project Last Checkpoint: 19 hours ago (unsaved changes) Logout
File Edit View Insert Cell Kernel Widgets Help Trusted Python 3
In [38]: from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import learning_curve

train_sizes = [1500, 2250, 2750, 3300, 3850, 4350, 4450]
tr_size, train_error, validation_error = learning_curve(
    estimator=Pipeline([
        ('bow', CountVectorizer(analyzer=text_process)), # strings to token integer counts
        ('tfidf', TfidfTransformer()), # integer counts to weighted TF-IDF scores
        ('classifier', MultinomialNB()), # train on TF-IDF vectors w/ Naive Bayes classifier
    ]),
    X=messages["message"], y=messages["label"], train_sizes=train_sizes, cv=5, scoring="neg_log_loss")

In [39]: train_error
Out[39]: array([[ -0.11031901, -0.10734751, -0.10949326, -0.10949326, -0.10949326],
 [-0.09724846, -0.0963419 , -0.09075512, -0.09185424, -0.09185424],
 [-0.09350201, -0.0928681 , -0.08740124, -0.08732897, -0.08732897],
 [-0.08643598, -0.08603058, -0.08146145, -0.08271708, -0.08271708],
 [-0.0822616 , -0.08231967, -0.07805742, -0.0808689 , -0.0810095 ],
 [-0.07809346, -0.078343 , -0.07453198, -0.07713932, -0.07734648],
 [-0.07699568, -0.07714761, -0.07345951, -0.07608351, -0.07635615]])

In [40]: validation_error
Out[40]: array([[ -0.15458714, -0.17166146, -0.16954966, -0.16299776, -0.15579178],
 [-0.13443863, -0.15070469, -0.15743567, -0.1458508 , -0.13692505],
 [-0.12394107, -0.1410872 , -0.14955551, -0.13650678, -0.12887588],
 [-0.11109497, -0.12847037, -0.13759886, -0.12932961, -0.12214693],
 [-0.10663434, -0.1195155 , -0.13323066, -0.12302478, -0.11556999],
 [-0.10312602, -0.11685917, -0.1297066 , -0.1176689 , -0.10854836],
 [-0.10252792, -0.1170407 , -0.12932392, -0.11725868, -0.10843111]])

In [42]: train_error_mean = -train_error.mean(axis = 1)
validation_error_mean = -validation_error.mean(axis = 1)
print('Mean training error\n', pd.Series(train_error_mean, index = train_sizes))
print('\n', '-' * 20) # separator
print('\nMean validation error\n', pd.Series(validation_error_mean, index = train_sizes))

Mean training error
1500    0.109229
2250    0.093611
2750    0.089686
3300    0.083872
3850    0.080903
4350    0.077091
4450    0.076008
dtype: float64

-----

Mean validation error
1500    0.162918
2250    0.145071
2750    0.135993
3300    0.125728
3850    0.119595
4350    0.115182
4450    0.114916
dtype: float64
```

Plotting Accuracy Curve:

```
In [50]: #Plotting accuracy curve
import matplotlib.pyplot as plt


plt.style.use('seaborn')
plt.plot(train_sizes, train_scores_mean, label = 'Training accuracy')
plt.plot(train_sizes, validation_scores_mean, label = 'Validation accuracy')
plt.ylabel('Accuracy', fontsize = 14)
plt.xlabel('Training set size', fontsize = 14)
plt.title('Learning curve for the naive bayes classification model', fontsize = 18, y = 1.03)
plt.legend()
plt.ylim(0.9,1)
```

Plotting Log Loss Curve:

```
In [46]: #Plotting loss curve
import matplotlib.pyplot as plt

plt.style.use('seaborn')
plt.plot(train_sizes, train_error_mean, label = 'Training error')
plt.plot(train_sizes, validation_error_mean, label = 'Validation error')
plt.ylabel('Log Loss', fontsize = 14)
plt.xlabel('Training set size', fontsize = 14)
plt.title('Log Loss curve for the naive bayes classification model', fontsize = 18, y = 1.03)
plt.legend()
plt.ylim(0,0.2)
```

Accuracy Graph:



Loss Graph:



Training and fitting the final model (Pipeline) on the cleaned dataset (With the best hyperparameters and the most optimal train-test split found above).

```
In [38]: from sklearn.model_selection import train_test_split

msg_train, msg_test, label_train, label_test = \
train_test_split(messages['message'], messages['label'], test_size=0.2)

print(len(msg_train), len(msg_test), len(msg_train) + len(msg_test))
4457 1115 5572
```

```
In [39]: from sklearn.pipeline import Pipeline

pipeline = Pipeline([
    ('bow', CountVectorizer(analyzer=text_process)), # strings to token integer counts
    ('tfidf', TfidfTransformer()), # integer counts to weighted TF-IDF scores
    ('classifier', MultinomialNB()), # train on TF-IDF vectors w/ Naive Bayes classifier
])
```

```
In [40]: pipeline.fit(msg_train, label_train)

Out[40]: Pipeline(steps=[('bow',
                          CountVectorizer(analyzer=<function text_process at 0x7fc11dc8f9d0>)),
                          ('tfidf', TfidfTransformer()),
                          ('classifier', MultinomialNB())])
```

```
In [41]: predictions = pipeline.predict(msg_test)

In [42]: print(classification_report(predictions, label_test))
```

	precision	recall	f1-score	support
ham	1.00	0.96	0.98	995
spam	0.76	0.99	0.86	120
accuracy			0.97	1115
macro avg	0.88	0.98	0.92	1115
weighted avg	0.97	0.97	0.97	1115

Methodology:

- Import dataset using pandas
- Exploratory data analysis to check how balanced the dataset actually is (Group by label).
- Add additional length column as a new text attribute for the final MultinomialNB ML model.
- Text pre-processing by removing punctuations, stopwords, etc.
- Tokenization, Normalization followed by Vectorization using CountVectorizer to create a Bag of Words representation of the data.
- Pass through TF-IDF transformer to get the new updated representation.
- Train test split the dataset using Sklearn Model Selection package.
- Create an Sklearn data pipeline and perform the above mentioned steps.

```

from sklearn.pipeline import Pipeline

pipeline = Pipeline([
    ('bow', CountVectorizer(analyzer=text_process)), # strings to token integer counts
    ('tfidf', TfidfTransformer()), # integer counts to weighted TF-IDF scores
    ('classifier', MultinomialNB()), # train on TF-IDF vectors w/ Naive Bayes classifier
])

```

- Fit the pipeline on the cleaned train dataset and transform the cleaned output dataset using the same trained pipeline.
- Check the overall final classification report.

```
In [60]: predictions = pipeline.predict(msg_test)
```

```
In [61]: print(classification_report(predictions, label_test))
```

	precision	recall	f1-score	support
ham	1.00	0.96	0.98	1001
spam	0.75	1.00	0.85	114
avg / total	0.97	0.97	0.97	1115

Result Metrics:

Final Accuracy: 97%

Final Recall: 97%

Final F1-Score: 97%

Final Support: 1115

Conclusion:

The final model achieves an overall test accuracy of 97 % and a recall of 97% which is really good considering the naturality of the real world dataset that has been considered.

References:

[1] <https://www.nltk.org/api/nltk.html>

[2] https://en.wikipedia.org/wiki/Natural_language_processing

[3] <https://www.ibm.com/cloud/learn/natural-language-processing>

[4] Deep Learning With Python. Manning Publications Co., 3 Lewis Street Greenwich, CT, United States

[5] https://www.sas.com/en_in/insights/analytics/what-is-natural-language-processing-nlp.html